Wei Cheng Chiu

LLM Serving & GPU Systems Engineer | On-prem LLM Deployment

waynehacking8@gmail.com | linkedin.com/in/wei-cheng-chiu | github.com/waynehacking8
wayne.is-a.dev | prs.wayne.is-a.dev

SUMMARY

Software engineer on **LLM infrastructure** at Taiwan AILabs: LLM inference, GPU performance, and on-prem GenAI deployment on NVIDIA GPUs. Upstream contributor with **22 merged/landed patches** across FlashInfer, vLLM, LMDeploy, Dynamo, PyTorch and torchao, including consumer-Blackwell (**SM120**) fixes.

TECHNICAL SKILLS

LLM Serving & GPU	vLLM, SGLang, TensorRT-LLM, Triton, FlashInfer, NCCL, NVFP4; CUDA, CUTLASS, Nsight
GenAI & Agents	Multi-agent orchestration (LangGraph), hybrid RAG, guardrails, tool use, MCP
Infrastructure	Linux, Kubernetes, Docker; H100, DGX Spark (GB10), Blackwell workstations
Languages	Python, C/C++, TypeScript, Rust; PyTorch, FastAPI

EXPERIENCE

Taiwan AILabs Jul 2026 – Present
Software Engineer, LLM Infrastructure *Taipei, Taiwan*

- Deploy and operate on-prem LLM systems (**FedGPT**) on Kubernetes and NVIDIA GPUs, from bare metal to handover

SYNCROBOTIC Sep 2025 – Jun 2026
Machine Learning Engineer *New Taipei City, Taiwan*

- **Sole developer** of an enterprise multi-agent RAG platform on DGX hardware, shipped to **two customers**

Advantech Jun 2025 – Aug 2025
AI Agent Engineer (Intern) *Taipei, Taiwan*

- Built tool-using coding agents for code review, test generation, and static analysis

TSMC (Taiwan Semiconductor Manufacturing Co.) May 2021 – Oct 2022
Facility Engineer (R&D Alternative Military Service) *Tainan, Taiwan*

- Supported R&D fab facility operations

OPEN SOURCE

Bug fixes across the LLM-inference stack 22 merged/landed
FlashInfer (9) · vLLM (5) · LMDeploy (3) · Dynamo (2) · PyTorch · torchao · compressed-tensors prs.wayne.is-a.dev

- Merged: FlashInfer stream-safe top-k, SM120 and NCCL fixes. In review: SM120 / NVFP4 in SGLang and vLLM

RESEARCH & PUBLICATIONS

SelGrad: Selective Gradient Projection for Efficient Safety Alignment Against Harmful Fine-Tuning, first author. *IEEE TDSC*, under review, 2026.

NTUST research in LLM security and privacy-preserving ML (federated learning, differential privacy, MPC).

EDUCATION

National Taiwan University of Science and Technology (NTUST) Aug 2024 – Apr 2026
M.S. in Computer Science and Information Engineering — GPA 4.09 *Taipei, Taiwan*

National Taiwan University (NTU) Sep 2014 – Jun 2020
M.S. in Structural Engineering (2020) | B.S. in Civil Engineering (2018) *Taipei, Taiwan*

CERTIFICATIONS & AWARDS

NVIDIA DLI — Accelerated Computing with CUDA C/C++ and CUDA Python (5 certificates) · NICS Privacy-Enhancing Technologies (PETs) Competition — Rank 3/11 (DP synthetic data via MP-SPDZ) · International ICT Innovative Services Awards, 2024